

workrr **one** — How It's Built & Deployed

A private business-intelligence platform that runs entirely inside your own Cloudflare account.

workrr one connects your business systems into a private, entity-resolved, time-aware **model of your business** (the “catalog”), and serves decision apps on top of it. It is built on Cloudflare's developer platform and **deployed 100% within your own Cloudflare account** — no workrr-hosted servers, no third-party model calls, no data egress.

Deployment model — your account, always


Every workrr one component is provisioned inside **your** Cloudflare account. workrr deploys and operates it there under scoped, auditable, revocable access — like a forward-deployed team with keys to one room, not a vendor holding your data.

- Compute, storage, vectors, and inference all live in your tenant.
- Your data never transits workrr infrastructure and is never sent to a model vendor.
- Nothing is shared across customers; there is no multi-tenant pool.

The stack (Cloudflare primitives)

LAYER	CLOUDFLARE SERVICE	ROLE
Connectors	Workers, Browser Rendering, Hyperdrive, Queues	Pull from M365/Google, CRM/ERP, SQL, files
Documents	R2	Source blobs (no egress fees)
Semantic layer	Workers AI → Vectorize	Embeddings + search / retrieval
Extraction	Workers AI / AI Gateway	Turn content into structured facts
Structured layer	D1 (SQLite)	Time-aware, sourced assertions & beliefs
Entity resolution / state	Durable Objects	Unify records; hot state; live updates
Pipeline	Workflows + Queues	Durable ingest → embed → extract → resolve
Decision apps	Workers + Analytics Engine	Blended metrics & always-live views
Access & security	Cloudflare Access / Zero Trust, WAF	Identity, governed & audited access

Data flow



Privacy & security

- ✓ **In-account:** all data and compute stay in your Cloudflare tenant.
- ✓ **No model-vendor egress:** inference runs on Workers AI (or your AI Gateway); prompts and data are not sent to OpenAI/Anthropic/etc.
- ✓ **No training on your data,** ever; no cross-customer pooling.
- ✓ **Zero Trust** identity for every user; workrr operator access is scoped, logged, and revocable.
- ✓ **Data localization** controls keep data in-region where required.
- ✓ **Sourced answers:** every result traces to the record it came from — auditable, not a black box.

Why this beats plain RAG or BI

The catalog has two planes

Semantic + structured, together.

- **Semantic:** search & ask across all content.
- **Structured:** entity-resolved, time-aware facts with provenance.

What that enables

- Joins the same customer/deal/dollar across systems.
- Answers “what was true when,” not just “what's a match.”
- Drives decisions, not just document snippets or static charts.

What a pilot looks like

- **Weeks, not quarters.** Connect 2–3 systems, stand up one decision app, validate on real data.
- **Scoped & low-risk:** runs in your account; remove access and it's gone.
- **Serverless economics:** scales with usage, no idle infrastructure, R2 with no egress fees.

What we need from you

- Scoped access to a Cloudflare account (yours — new or existing).
- Read access to the 2–3 source systems for the first decision app.
- A business sponsor who owns the reporting ritual we're replacing.